

Tab 1

THE COMPLETE GUIDE

Cut your Claude token usage by 80%

15 simple habits that save tokens, reduce costs, and make every Claude session more productive. No hacks. Just better workflows.

*Simple changes, **massive results**.*

Why tokens matter more than you think

Claude doesn't count messages. It counts tokens. Every word, every space, every piece of punctuation.

Here's the part most people miss: Claude re-reads your **entire conversation** every time you hit send. Message 1 costs 500 tokens. By message 15, that same simple question might cost 15,000 to 25,000 tokens because it's carrying the full weight of every previous message and response.

So a 30-message chat isn't 30x the cost of one message. It's closer to 200x or more. That's why these habits matter. Small changes in how you work with Claude add up to massive savings.

01. Never let a chat go past 20 messages

Longer chat = more tokens. Every new message forces Claude to re-read everything that came before it. By message 20, you're paying for a massive context window even if your actual question is tiny.

After 15 to 20 messages, ask Claude for a summary of the full conversation. Copy it. Open a new chat. Paste the summary. Keep going from there.

You get the same continuity with a fraction of the token cost.

Pro move

Type “Summarize this entire chat in a way I can paste into a new conversation to continue where we left off.” Claude will give you a clean handoff every time.

02. Tell Claude how long to respond

Claude defaults to long, thorough answers full of filler. That’s great when you need depth. Most of the time, you don’t.

Add a length constraint to your prompt and you directly control how many output tokens Claude uses.

✗ COSTS MORE

“Explain how prompt caching works” Claude writes 400+ words with background, history, and caveats you didn’t ask for.

✓ COSTS LESS

“Explain prompt caching in 2 sentences” Claude gives you exactly what you need. Done.

Try these phrases

“Answer in 2 sentences.” / “Keep it under 50 words.” / “One paragraph max.” / “Just the code, no explanation.”

03. Use Projects to store context

Stop pasting your company info, brand guidelines, and docs into every single chat. That’s tokens wasted on repeat.

Put all your reference material inside a Claude Project instead. Claude reads it automatically in every conversation within that project. You write it once, you save tokens forever.

Think of Projects as Claude’s long-term memory for a specific client or workflow. Your brand voice doc, your audience notes, your style guide, all of it lives there and doesn’t need to be re-sent.

How to set it up

Go to your Claude sidebar, click “New project,” give it a name, and upload your reference files. Every new chat inside that project has access to those files automatically.

04. Don’t always let Claude ask questions

By default, Claude loves to ask clarifying questions before it does anything. That feels polite but it costs you an extra round-trip. Two messages instead of one. Double the tokens.

When your prompt has enough context, add this line: “Don’t ask questions. Make your best choice based on the context I gave you.”

Fewer turns = fewer tokens. Simple math.

✗ 2 TURNS, 2X TOKENS

You: “Write a blog intro” Claude: “What topic? What audience?” You: “SaaS founders, about pricing” Claude: [finally writes it]

✓ 1 TURN, 1X TOKENS

You: “Write a blog intro about SaaS pricing for startup founders. 3 sentences. Don’t ask questions.” Claude: [writes it immediately]

05. Edit and regenerate instead of following up

Didn’t like the output? The instinct is to send a follow-up: “Make it shorter” or “Change the tone.” That adds more history and more tokens to process.

Instead, go back to your original prompt, edit it directly, and hit regenerate. The old reply disappears. Claude starts fresh with just one clean prompt. No accumulated history. No token bloat.

This single habit can cut your per-task token cost in half.

When to follow up instead

If you’re building on something (like iterating on code where each step depends on the last), follow-ups make sense. But for standalone tasks like writing a caption or email, always edit and regenerate.

06. Make Claude remember you

Stop re-explaining your role, your tone, your audience, and your preferred format in every chat. That's hundreds of wasted tokens per conversation.

Go to Settings, then General, then "Instructions to Claude." Write your details there once. Claude carries this into every conversation automatically.

What to include: *Your name and role, what you work on, your audience, your writing style (with an example), formatting preferences, and words to avoid.*

07. Stop regenerating blindly

Don't hit regenerate 5 times hoping for a better answer. Each one costs the same number of tokens. You're gambling, not improving.

If the first output missed the mark, the problem isn't luck. It's the prompt. Change what you asked for. Be more specific about what was wrong. Then regenerate once.

One good prompt beats five random rolls every time.

8 MORE HABITS THAT SAVE TOKENS

Beyond the basics. These are the habits that separate casual users from power users.

08. Batch your requests into one message

Sending five small follow-ups ("change this," "now fix that," "also adjust this") is one of the most expensive patterns. Each message forces Claude to reprocess the entire conversation from scratch.

Collect all your changes and send them in one shot.

✗ 5 MESSAGES = 5X REPROCESSING

"Make the headline shorter" "Change CTA to Start free" "Add a testimonial" "Make it more casual" "Fix the typo in line 2"

✓ 1 MESSAGE = 1X PROCESSING

"Update the landing page copy: shorter headline, change CTA to Start free, add a testimonial section, more casual tone, and fix the typo in line 2."

09. Use XML tags to organize your prompts

When you mix instructions with data in one big paragraph, Claude sometimes gets confused. It spends tokens figuring out what's the task and what's the content. That confusion leads to longer, less accurate responses.

Wrap your data in simple tags. Keep your instructions outside them. Claude processes this much more efficiently.

✗ MESSY PROMPT Here's my article about AI tools for startups, it's about 2000 words and I need you to rewrite the intro to be punchier and also check for any factual errors...	✓ STRUCTURED PROMPT <task>Rewrite the intro to be punchier. Flag any factual errors.</task><article> [your article text here] </article>
---	--

Clean input = clean output = fewer wasted tokens.

10. Pick the right model for the job

Not every task needs the most powerful model. Using Opus for a quick formatting fix is like taking a helicopter to the grocery store.

Match the model to the task. Simple edits, formatting, quick Q&A? Sonnet or Haiku will handle it faster and cheaper. Save the heavy models for deep research, complex writing, and multi-step reasoning.

Quick rule If the task takes you less than 2 minutes to explain, a lighter model is probably enough. If it needs deep thinking or nuanced judgment, go with the bigger one.

11. Cut the polite filler from your prompts

It feels natural to write "Could you please help me with..." or "I was wondering if you might be able to..." But those extra words cost tokens and add zero value to the output.

Claude doesn't need pleasantries. It needs clear instructions.

✗ 32 TOKENS OF POLITENESS

Hi Claude! I was hoping you could help me out. Could you please write a short Instagram caption about our new product launch? Thank you so much!

✓ 14 TOKENS, SAME RESULT

Write a short Instagram caption about our new product launch.

This doesn't mean you need to be rude. Just direct. Think command line, not conversation.

12. Trim your input before pasting

Claude processes everything you send, even if most of it isn't relevant. Pasting a 5,000-word article when you only need feedback on the intro? That's 4,500 words of wasted input tokens.

Before you paste, ask yourself: does Claude need all of this? If not, cut it down to only the relevant parts. Less input = less token usage, always.

Example

Instead of pasting your entire website copy for a headline review, just paste the headline and the first paragraph. Give Claude only what it needs to do the job.

13. Give one example instead of explaining

A 200-word description of your desired tone costs more tokens than just showing Claude what you want. One good example replaces paragraphs of instructions.

✗ LONG EXPLANATION

I want the tone to be casual but not too casual, professional but approachable, with short sentences and a conversational feel, kind of like talking to a friend but still smart and credible...

✓ ONE EXAMPLE

Write in this tone: "Most founders overthink their first landing page. Say what it does. Say who it's for. Add a button. Ship it."

Show, don't tell. It's cheaper and it works better.

14. Spread your sessions across the day

Claude's usage limits work on a rolling window, not a midnight reset. If you dump all your heavy tasks into one 3-hour session, you burn through your budget fast and leave the rest of your day empty.

Split your work into blocks. Research in the morning. Writing after lunch. Review in the evening. By the time you start your second session, the first one is already rolling off your limit.

Same total work, but your usage stays under the ceiling.

15. Turn off features you don't need

Extended thinking, web search, and other features are powerful. They also multiply your token usage significantly before Claude even starts writing your answer.

If you're brainstorming product ideas, you don't need web search. If you're writing a caption, you don't need extended thinking. Turn off what you're not actively using.

Every feature you leave on is a feature that's eating tokens in the background.

Where to toggle

You can turn features on and off right from the chat input area. Web search, extended thinking, and code execution all have toggles. Default to off, turn on only when needed.

Your action checklist

Pick 3 habits to start with. Build from there. You'll see the difference in your very first session.

- ☐ Cap chats at 20 messages and use the summary-paste workflow
- ☐ Add length constraints to every prompt ("2 sentences," "under 50 words")
- ☐ Set up a Claude Project with your brand docs and guidelines
- ☐ Add "Don't ask questions" to prompts when context is clear
- ☐ Edit and regenerate instead of sending follow-up corrections
- ☐ Fill out Settings > General > Instructions to Claude
- ☐ Fix the prompt before hitting regenerate
- ☐ Batch multiple edits into a single message
- ☐ Use XML tags to separate instructions from data

- ☐ Use lighter models (Sonnet/Haiku) for simple tasks
- ☐ Drop polite filler from prompts and write direct commands
- ☐ Trim large inputs to only the relevant sections
- ☐ Show one example instead of writing long tone descriptions
- ☐ Spread heavy Claude sessions across the day
- ☐ Turn off extended thinking and web search when not needed

Start with just 3 habits.

You don't need all 15 at once. Pick the three that match how you work right now. Build from there. You'll see the difference in your very first session.

Simple changes, massive results.